

AI 1 Project

Customer Relationship Prediction

Team: Flip-A-Coin

Members:

1. Shambhavi Aggarwal
2. Lakshay Chawla
3. Devansh Shrestha
4. Rommel Jalasutram

The following cell outlines the methodology we followed:

1. Feature Selection
2. Train against different models
3. Results

Discuss more about this

In []:

```
## Importing the required libraries
## For dealing with data.

import numpy as np
import pandas as pd

## For preparation of data - This includes cleaning the data, feature selection, splitting
data into training - testing.

from sklearn.impute import SimpleImputer
from sklearn.preprocessing import StandardScaler

from sklearn.feature_selection import mutual_info_classif
from sklearn.ensemble import RandomForestClassifier

from sklearn.model_selection import train_test_split

# Models to be used.

from sklearn.model_selection import cross_validate
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from imblearn.under_sampling import RandomUnderSampler
from imblearn.over_sampling import SMOTE
from sklearn.ensemble import AdaBoostClassifier
from sklearn.model_selection import GridSearchCV

# For interpretation

import matplotlib.pyplot as plt
import seaborn as sns
from prettytable import PrettyTable

# Comparison Metrics

from sklearn.inspection import permutation_importance
from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay, roc_curve, f1_score,
roc_auc_score, accuracy_score, precision_score, recall_score
```

In []:

```

## Read the training data
## Dropping the 1st column as it has newline character
df = pd.read_csv("orange_small_train.data", sep = "\t", lineterminator="\r")
df = df.drop(['Var1'], axis = 1)

y_train_upselling = pd.read_csv("/content/orange_small_train_upselling.labels", sep = "\t", lineterminator="\n", names= ["upselling"])
y_train_appetency = pd.read_csv("/content/orange_small_train_appetency.labels", sep = "\t", lineterminator="\n", names= ["appetency"])
y_train_churn = pd.read_csv("/content/orange_small_train_churn.labels", sep = "\t", lineterminator="\n", names= ["churn"])

```

In []:

```

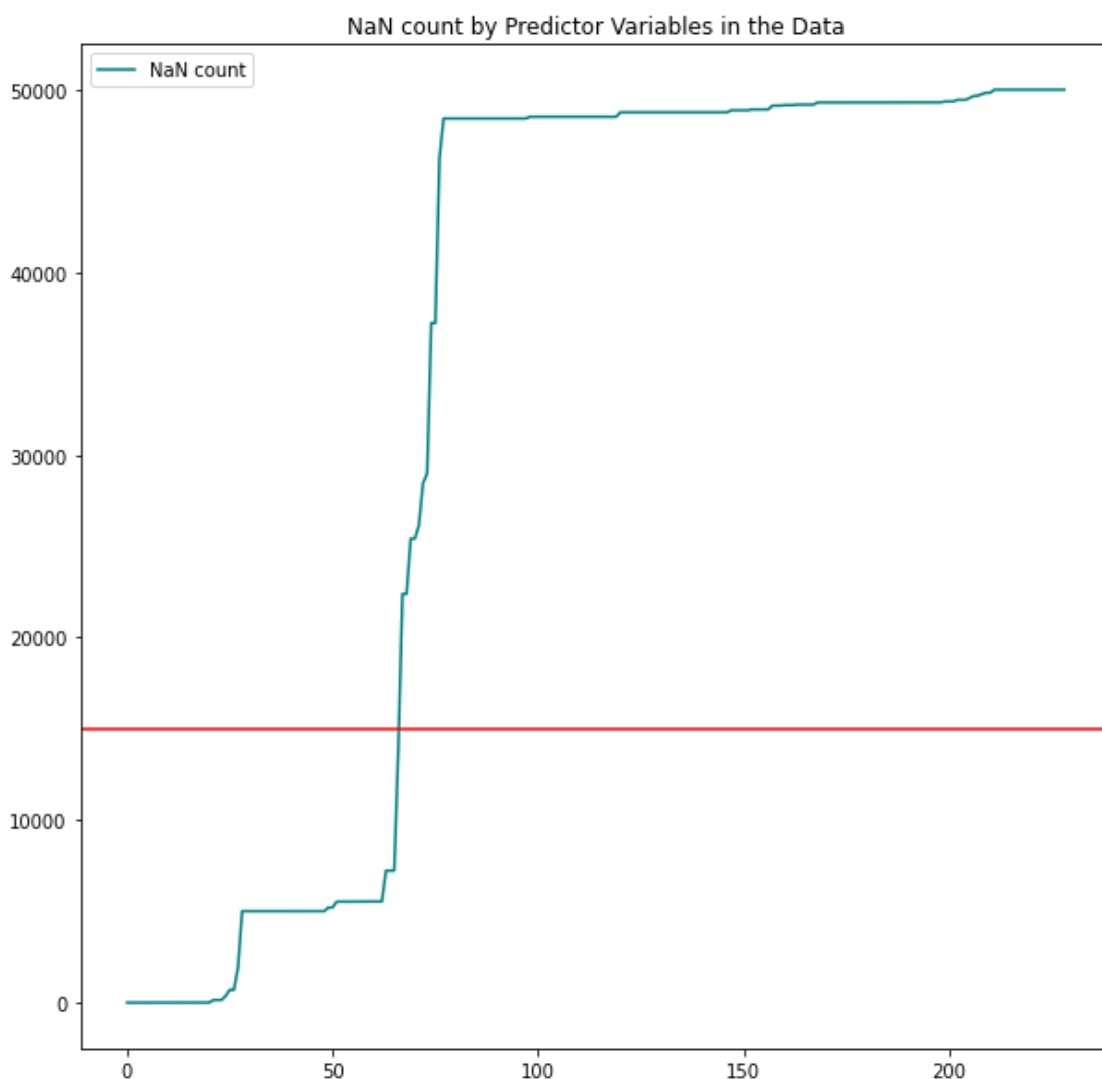
## Graph showing number of Nan values in each column
nan_count = []
for col in df.columns:
    nan_count.append(df[col].isna().sum())

nan_count.sort()
nan_c = pd.Series(nan_count)
plt.rcParams["figure.figsize"] = (10,10)
nan_c.plot(kind= "line", color = "teal", label = "NaN count")
plt.axhline(15000, c='r')
plt.legend(loc = 'upper left')
plt.title('NaN count by Predictor Variables in the Data')

```

Out[]:

Text(0.5, 1.0, 'NaN count by Predictor Variables in the Data')



In []:

```

## Plot showing the unique categories by categorical variables
num_unique_categories = []

```

```


for col in df.columns[191:]:
    num_unique_categories.append(len(df[col].unique()))

num_unique_categories.sort()
cat_unique = pd.Series(num_unique_categories)
plt.rcParams["figure.figsize"] = (10,10)
cat_unique.plot(kind= "line", color = "teal", label = "Category count")
plt.yscale('log')
plt.axhline(9, c='r')
plt.xlabel('Categories')
plt.ylabel('Unique Category Count (log scale)')
plt.legend(loc = 'upper left')
plt.title('Unique Category Count for Category Variables in the Data')

```

Out[]:

Text(0.5, 1.0, 'Unique Category Count for Category Variables in the Data')



In []:

```

## Several columns have NaN values
## For our purpose, a useful column is one that has fewer than 15000 NaNs
## or 30% NaNs.
def get_useful_columns(data_type):
    useful_cols = []

    for col in df.columns:
        nan_count = df[col].isna().sum()
        if df[col].dtypes == data_type and nan_count < 15000:
            if data_type == 'float64':
                useful_cols.append(col)
            elif data_type == 'object' and len(df[col].unique()) < 9:
                useful_cols.append(col)

```

```
return useful_cols
```

In []:

```
## Impute numerical data as several columns have NaN values
## The default metric is median.
def get_clean_numerical_data(num_cols_list, data_df, metric="median"):
    df_num = data_df[num_cols_list].copy()
    if metric == "median":
        df_num.fillna(df_num.median(),inplace = True)
    elif metric == "mean":
        df_num.fillna(df_num.mean(),inplace = True)

    ## Standardize the values as there are outliers
    scalar = StandardScaler()
    df_std = pd.DataFrame(scalar.fit_transform(df_num),columns = num_cols_list)
    return df_std
```

In []:

```
## Preprocessing the categorical columns
useful_cat_cols = get_useful_columns('object')
df_categorical = df[useful_cat_cols].copy()

## 1. Imputing missing values
simple_imputer = SimpleImputer(strategy = "most_frequent")
df_categorical = pd.DataFrame(simple_imputer.fit_transform(df_categorical), columns = use
ful_cat_cols)
```

In []:

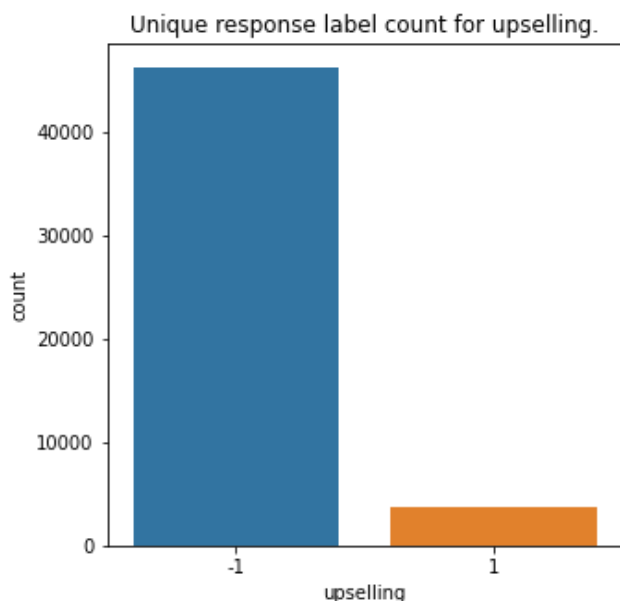
```
## Preprocessing the numerical columns
useful_num_cols = get_useful_columns('float64')
df_numerical = get_clean_numerical_data(useful_num_cols, df)
```

In []:

```
## Unique response label count for upselling.
plt.rcParams["figure.figsize"] = (5,5)
sns.countplot(x= "upselling", data=y_train_upselling)
plt.title('Unique response label count for upselling.')
```

Out[]:

Text(0.5, 1.0, 'Unique response label count for upselling.')



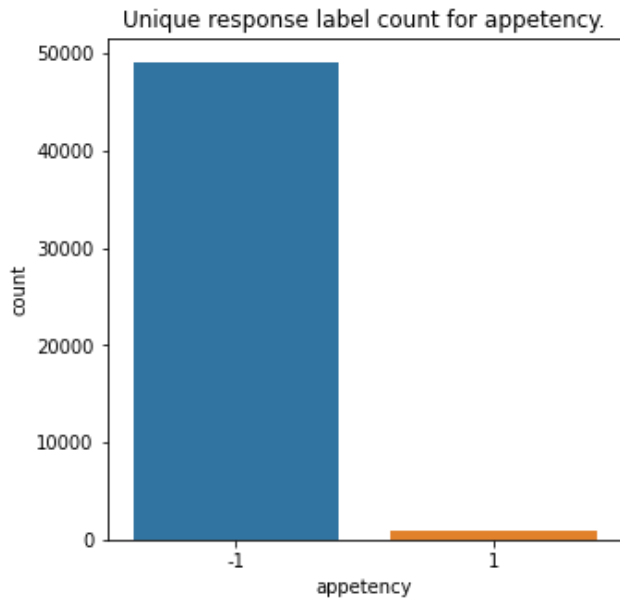
In []:

```
## Unique response label count for appetency.
plt.rcParams["figure.figsize"] = (5,5)
```

```
sns.countplot(x= "appetency", data=y_train_appetency)
plt.title('Unique response label count for appetency.')
```

Out[]:

Text(0.5, 1.0, 'Unique response label count for appetency.')

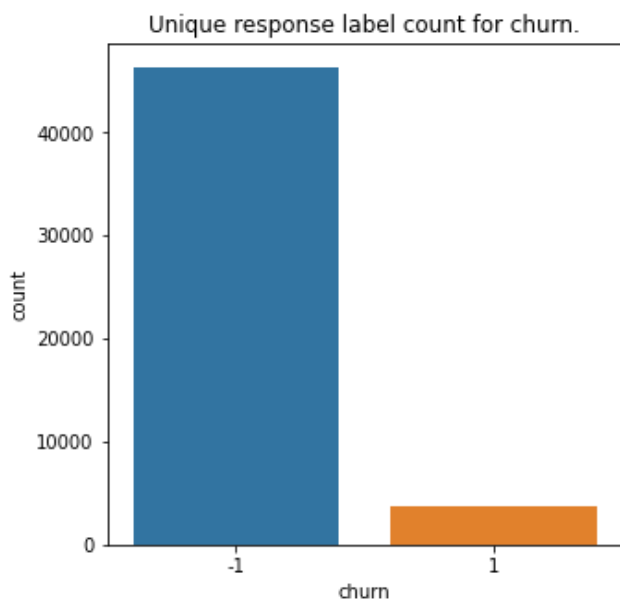


In []:

```
## Unique response label count for churn.
plt.rcParams["figure.figsize"] = (5,5)
sns.countplot(x= "churn", data=y_train_churn)
plt.title('Unique response label count for churn.')
```

Out[]:

Text(0.5, 1.0, 'Unique response label count for churn.')



In []:

```
## Putting it all together
## Creating the data frame we will be working with
df_train = pd.concat([df_numerical, df_categorical], axis = 1)
df_train = df_train[:-1]
df_train = pd.get_dummies(df_train, columns = useful_cat_cols)
```

Methods

In []:

```

## Intersection of lists having Important features from different methods.
def get_feature_intersection(lst1, lst2):
    ans = [value for value in lst1 if value in lst2]
    return ans

```

In []:

```

## Splitting the training data as there is no corresponding label file
## for orange_small_test.data
def get_train_test_split(df, y, random_state = 42):
    x_train, x_test, y_train, y_test = train_test_split(df,
                                                         y,
                                                         train_size = 0.80,
                                                         random_state = random_state,
                                                         stratify = y)

    return x_train, x_test, y_train, y_test

```

In []:

```

## Logistic Regression
def get_logistic_regression(X, y):
    random_state = 42
    c_list = [1e-5, 1e-4, 1e-3, 1e-2, 1e-1, 1, 1e1, 1e2, 1e3, 1e4, 1e5]
    k = 5
    lreg = LogisticRegression(class_weight= 'balanced', solver = 'lbfgs', random_state =
random_state, n_jobs=-1, max_iter = 5000)
    param_grid = {'C': c_list}
    scoring = {'AUC': 'roc_auc'}
    grid_search = GridSearchCV(lreg,
                               param_grid = param_grid,
                               scoring = scoring,
                               refit = 'AUC',
                               return_train_score = True,
                               cv = k,
                               n_jobs = -1)

    results = grid_search.fit(X, y)
    print(results.best_estimator_.get_params())

    ## Get the best estimator
    best_rf = results.best_estimator_
    return best_rf, grid_search

```

In []:

```

## Decision Tree
def get_decision_tree(x,y):
    depth_list = [3,6,9,12,15]
    split=[3,4,5,6]
    r_state = 42
    k=5
    scoring = {'AUC': 'roc_auc'}
    param_grid={'max_depth':depth_list,'min_samples_split':split}

    model=DecisionTreeClassifier(random_state =r_state)

    grid_search = GridSearchCV(model,
                               param_grid = param_grid,
                               scoring = scoring,
                               refit = 'AUC',
                               return_train_score=True,
                               cv=k,
                               n_jobs=-1)

    result = grid_search.fit(x,y)
    print(result.best_estimator_.get_params())
    best_rf = result.best_estimator_
    return best_rf, grid_search

```

In []:

```

## Vanilla Random Forest
## Return the best RF after grid search
def get_rf_vanilla(X, y):
    seed = 0
    max_depths = [10, 20]
    estimators = [20, 40, 60, 80, 100]

    param_grid = {'n_estimators': estimators, 'max_depth': max_depths}
    scoring = {'AUC': 'roc_auc'}

    random_forest = RandomForestClassifier(n_jobs = -1,
                                          n_estimators = 0,
                                          oob_score=True,
                                          max_features = 'sqrt',
                                          random_state = seed)

    grid_search = GridSearchCV(random_forest,
                               param_grid = param_grid,
                               scoring = scoring,
                               refit = 'AUC',
                               return_train_score = True,
                               n_jobs = -1)

    results = grid_search.fit(X, y)
    print(results.best_estimator_.get_params())

    ## Get the best estimator
    best_rf = results.best_estimator_
    return best_rf, grid_search

```

In []:

```

## Balancing the class imbalance in each bootstrap
## Using the parameters for best tree from grid search for vanilla random forest
def get_rf_balance(X, y):
    random_forest = RandomForestClassifier(n_jobs = -1,
                                          max_depth = 10,
                                          n_estimators = 100,
                                          oob_score=True,
                                          class_weight = 'balanced_subsample')

    random_forest.fit(X, y)
    return random_forest

## Refit the model with the best parameters, calculate the auc on the test data

```

In []:

```

## Balanced Random Forest with SMOTE
## Using the parameters for best tree from grid search for vanilla random forest
def get_rf_smote(X, y):
    sm = SMOTE(random_state = 2)
    X_res, y_res = sm.fit_resample(X, y)

    random_forest = RandomForestClassifier(n_jobs = -1,
                                          max_depth = 10,
                                          n_estimators = 100,
                                          oob_score=True,
                                          class_weight = 'balanced_subsample')

    random_forest.fit(X_res, y_res)
    return random_forest

## Refit the model with the best parameters, calculate the auc on the test data

```

In []:

```

## Balanced Random Forest with RandomUnderSampler.
## Using the parameters for best tree from grid search for vanilla random forest
def get_rf_downsampler(X, y):
    rf = RandomUnderSampler(random_state = 2)

```

```

X_res, y_res = rf.fit_resample(X, y)

random_forest = RandomForestClassifier(n_jobs = -1,
                                     max_depth = 10,
                                     n_estimators = 100,
                                     oob_score=True,
                                     class_weight = 'balanced_subsample')

random_forest.fit(X_res, y_res)
return random_forest

## Refit the model with the best parameters, calculate the auc on the test data

```

In []:

```

## AdaBoost
## Return the best estimator after grid search
def get_adaboost(X, y, refit = 'AUC'):
    boost = AdaBoostClassifier( base_estimator = DecisionTreeClassifier(max_depth = 1),
                               algorithm = 'SAMME', n_estimators=0)

    learning_rates = [1e-2, 1e-1, 1, 10]
    estimators = [20, 40, 60, 80, 100]
    scoring = {'AUC': 'roc_auc', 'PREC': 'precision', 'RECALL': 'recall'}

    param_grid = {'n_estimators': estimators, 'learning_rate': learning_rates}
    grid_search = GridSearchCV(boost,
                               param_grid = param_grid,
                               scoring = scoring,
                               refit = refit,
                               return_train_score = True,
                               n_jobs = -1)

    results = grid_search.fit(X, y)
    print(results.best_estimator_.get_params())

    ## Get the best estimator
    best_rf = results.best_estimator_
    return best_rf, grid_search

```

In []:

```

## Get all the comparison metrics.
def get_metrics(predictions, labels, probabilities):
    ## Calculating the metrics.
    metrics = {}
    metrics["accuracy"] = accuracy_score(labels, predictions)
    metrics["recall"] = recall_score(labels, predictions)
    metrics["precision"] = precision_score(labels, predictions)
    metrics["f1_score"] = f1_score(labels, predictions)
    metrics["auc_score"] = roc_auc_score(labels, probabilities[:, 1])
    return metrics

```

In []:

```

## Print the metrics for all the models in a pretty table.
def print_metrics(metrics_dict):
    pt = PrettyTable()
    pt.field_names = ["Model Name", "Accuracy", "Recall", "Precision", "F1 Score", "AUC score"]
    for i in metrics_dict.keys():
        pt.add_row([i, metrics_dict[i]['accuracy'], metrics_dict[i]['recall'], metrics_dict[i]['precision'], metrics_dict[i]['f1_score'], metrics_dict[i]['auc_score']])
    print(pt)

```

In []:

```

## Get the permutation importance
def get_permutation_importance(model, x, y, ff):
    ## Calculate the Permutation importance
    r = permutation_importance(model, x, y,
                              n_repeats=30,
                              random_state=0)

```



```

for i in r.importances_mean.argsort()[::-1]:
    if r.importances_mean[i] - 2 * r.importances_std[i] > 0:
        print(f"{ff[i]:<8}"
              f"{r.importances_mean[i]:.3f}"
              f" +/- {r.importances_std[i]:.3f}")

```

In []:

```

## Code to plot confusion matrix
def get_confusion_matrix(y_orig, y_pred, classes):
    plt.rcParams["figure.figsize"] = (10,10)
    conf_matrix = confusion_matrix(y_orig, y_pred)
    conf_matrix_plot = ConfusionMatrixDisplay(confusion_matrix = conf_matrix, display_labels = classes)
    conf_matrix_plot.plot()
    return conf_matrix

```

In []:

```

## Getting ROC-AUC curve
def get_roc_curve(model, testX, testy, name):
    plt.rcParams["figure.figsize"] = (10,10)
    y_probs = model.predict_proba(testX)[:,1]
    fpr, tpr, threshold = roc_curve(testy, y_probs)
    plt.plot(fpr, tpr, color = 'red')
    plt.plot([0, 1], [0, 1], 'r--', color = 'blue')
    plt.xlabel('False Positive Rate')
    plt.ylabel('True Positive Rate')
    plt.title(f'ROC curve for {name}.')
    plt.legend(loc="lower right")
    return fpr, tpr

```

Predicting Upselling

In []:

```

## Get the test train split with Upselling data
x_train, x_test, y_train, y_test = get_train_test_split(df_train, y_train_upselling)

```

Feature Selection

We have 81 columns. We try to find the most useful columns from these based on:

1. Information gain
2. Random Forest Feature Importance

Information gain

In []:

```

## Finding important features using information gain.
ig_importances = mutual_info_classif(x_train, y_train.values.ravel())
feat_importances = pd.Series(ig_importances, df_train.columns)

## Plot to show important features according to Information Gain.
plt.rcParams["figure.figsize"] = (25,60)
feat_importances.plot(kind= "barh", color = "teal")
plt.title('Plot to show important features according to Information Gain.')
plt.show()


```


Plot to show important features according to Information Gain.

Var227_vj_w8kB

Var227_nIGjg5B

Var227_nIGXDli





In []:

```
## Boxplot of important features according to Information Gain.
plt.rcParams["figure.figsize"] = (5,5)
sns.boxplot(feat_importances).set_title('Boxplot of important features according to Information Gain.')
```


/usr/local/lib/python3.7/dist-packages/seaborn/_decorators.py:43: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

FutureWarning

Out[]:

Text(0.5, 1.0, 'Boxplot of important features according to Information Gain.')

Boxplot of important features according to Information Gain.



In []:


```
## Removing 25% of least important features.
ft_ig = feat_importances[feat_importances >= feat_importances.quantile(0.25)].index.tolist()
```


Random Forest Feature Importance and Feature Selection

In []:

```
## Finding important features using Random Forest Feature Selection.
rf_feat_select = RandomForestClassifier(n_estimators=340)
rf_feat_select.fit(x_train, y_train.values.ravel())
rf_imp = rf_feat_select.feature_importances_
rf_importance = pd.Series(rf_imp, df_train.columns)

## Plot to show important features according to Random Forest Feature Selection.
plt.rcParams["figure.figsize"] = (25,60)
rf_importance.plot(kind= "barh", color = "teal")
plt.show()
```





In []:

```
## Boxplot of important features according to Random Forest Feature Selection.
plt.rcParams["figure.figsize"] = (5,5)
sns.boxplot(rf_importance)
```


/usr/local/lib/python3.7/dist-packages/seaborn/_decorators.py:43: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

FutureWarning

Out[]:

<matplotlib.axes._subplots.AxesSubplot at 0x7fe7ac5eae50>





In []:

```
## Removing 25% of least important features.
ft_rf = rf_importance[rf_importance >= rf_importance.quantile(0.25)].index.tolist()
```

In []:

```
## Using the intersection function to get the final list of important features.
final_features= get_feature_intersection(ft_rf, ft_ig)
```

Model

In []:

```
## Select only the important features in x_train and x_test.
x_train = pd.DataFrame(x_train, columns=df_train.columns)
x_test = pd.DataFrame(x_test, columns=df_train.columns)

x_train = x_train[final_features].values
x_test = x_test[final_features].values
```

In []:


```
## The dictionaries hold metrics, confusion matrix data, and ROC-AUC curve information
## for all the models
final_result_upselling = {}
conf_matrix_upselling = {}
roc_upselling = {}
```

In []:

```
## Linear Regression model
lreg, gs_lreg = get_logistic_regression(x_train, y_train.values.ravel())
predictions = lreg.predict(x_test)
probabilities = lreg.predict_proba(x_test)
final_result_upselling["logic_reg"] = get_metrics(predictions, y_test.values.ravel(), probabilities)

conf_matrix_upselling["logic_reg"] = get_confusion_matrix(y_test, predictions, lreg.classes_)

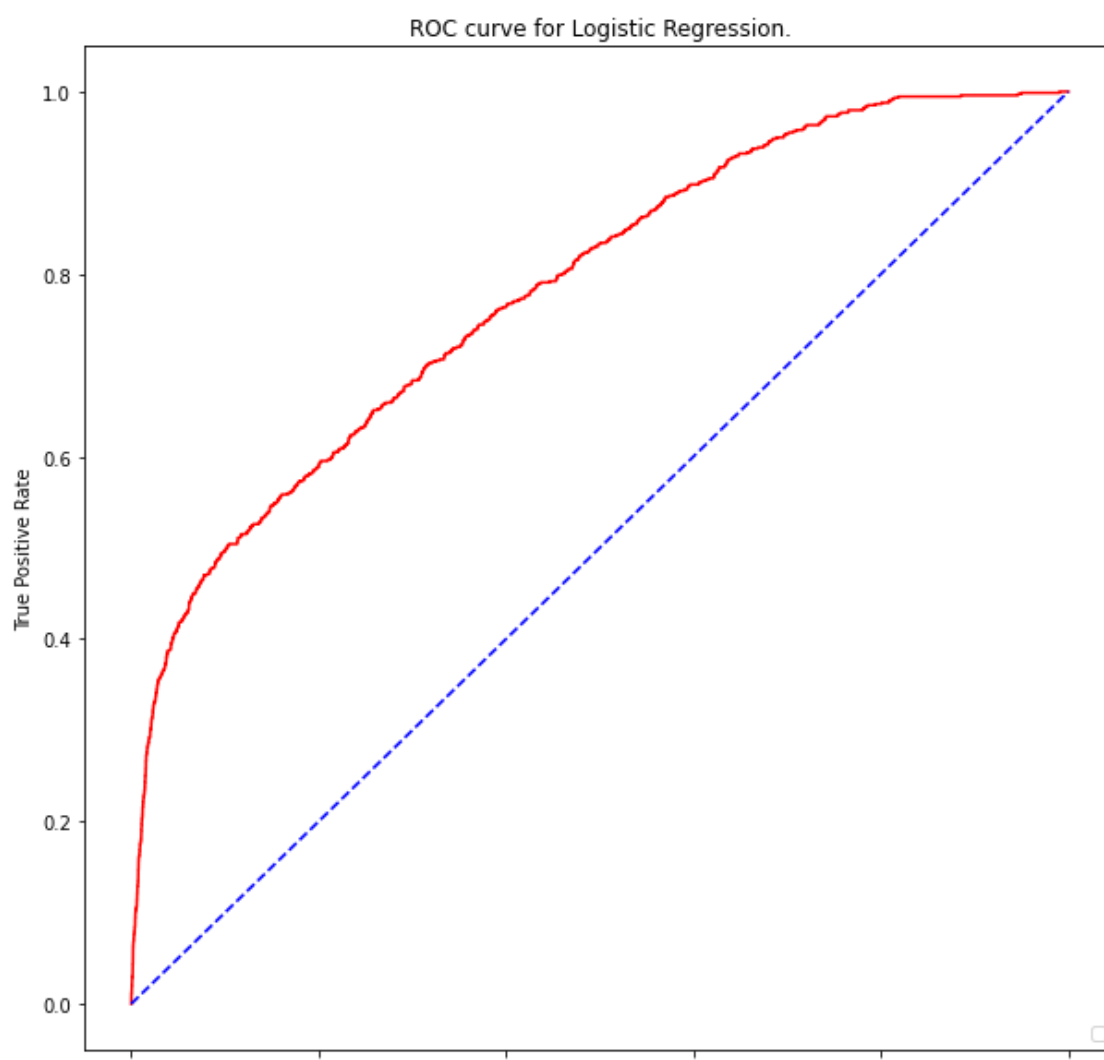
{'C': 10.0, 'class_weight': 'balanced', 'dual': False, 'fit_intercept': True, 'intercept_scaling': 1, 'l1_ratio': None, 'max_iter': 5000, 'multi_class': 'auto', 'n_jobs': -1, 'penalty': 'l2', 'random_state': 42, 'solver': 'lbfgs', 'tol': 0.0001, 'verbose': 0, 'warm_start': False}
```



In []:

```
## ROC Curve for Logistic Regression.
roc_upselling["logic_reg"] = get_roc_curve(lreg, x_test, y_test, "Logistic Regression")
```

No handles with labels found to put in legend.




In []:

```
## Decision Tree Model.
dst, gs_dst = get_decision_tree(x_train, y_train.values.ravel())
predictions = dst.predict(x_test)
probabilities = dst.predict_proba(x_test)
final_result_upselling["decision_tree"] = get_metrics(predictions, y_test.values.ravel(),
, probabilities)

conf_matrix_upselling["decision_tree"] = get_confusion_matrix(y_test, predictions, dst.c
lasses_)


{'ccp_alpha': 0.0, 'class_weight': None, 'criterion': 'gini', 'max_depth': 6, 'max_featur
es': None, 'max_leaf_nodes': None, 'min_impurity_decrease': 0.0, 'min_samples_leaf': 1, '
min_samples_split': 3, 'min_weight_fraction_leaf': 0.0, 'random_state': 42, 'splitter': '
best'}
```




In []:

```
## ROC Curve for Decision Tree.
roc_upselling["decision_tree"] = get_roc_curve(dst, x_test, y_test, "Decision Tree")
```

No handles with labels found to put in legend.






In []:

```
## Vanilla Random Forest Model
rfv, gs_rfv = get_rf_vanilla(x_train, y_train.values.ravel())
predictions = rfv.predict(x_test)
probabilities = rfv.predict_proba(x_test)
final_result_upselling["vanilla_random_forest"] = get_metrics(predictions, y_test.values
.ravel(), probabilities)

conf_matrix_upselling["vanilla_random_forest"] = get_confusion_matrix(y_test, predictions
, rfv.classes_)

{'bootstrap': True, 'ccp_alpha': 0.0, 'class_weight': None, 'criterion': 'gini', 'max_dep
th': 10, 'max_features': 'sqrt', 'max_leaf_nodes': None, 'max_samples': None, 'min_impuri
ty_decrease': 0.0, 'min_samples_leaf': 1, 'min_samples_split': 2, 'min_weight_fraction_le
af': 0.0, 'n_estimators': 100, 'n_jobs': -1, 'oob_score': True, 'random_state': 0, 'verbo
se': 0, 'warm_start': False}
```






In []:

```
## ROC Curve for vanilla random forest.
roc_upselling["vanilla_random_forest"] = get_roc_curve(rfv, x_test, y_test, "Vanilla Random Forest")
```

No handles with labels found to put in legend.




In []:

```
## Balanced Random Forest Model.
rf_balance = get_rf_balance(x_train, y_train.values.ravel())
predictions = rf_balance.predict(x_test)
probabilities = rf_balance.predict_proba(x_test)
final_result_upselling["rf_balanced"] = get_metrics(predictions, y_test.values.ravel(), probabilities)

conf_matrix_upselling["rf_balanced"] = get_confusion_matrix(y_test, predictions, rf_balance.classes_)
```






In []:

```
## ROC Curve for balanced random forest.
roc_upselling["rf_balanced"] = get_roc_curve(rf_balance, x_test, y_test, "Balanced Random Forest")
```

No handles with labels found to put in legend.




In []:

```

## Random Forest model with smote upsampling.
rf_smote = get_rf_smote(x_train, y_train.values.ravel())
predictions = rf_smote.predict(x_test)
probabilities = rf_smote.predict_proba(x_test)
final_result_upselling["rf_smote"] = get_metrics(predictions, y_test.values.ravel(), probabilities)

conf_matrix_upselling["rf_smote"] = get_confusion_matrix(y_test, predictions, rf_smote.classes_)

```




In []:


```

## ROC Curve for Random Forest model with smote upsampling.
roc_upselling["rf_smote"] = get_roc_curve(rf_smote, x_test, y_test, "Random Forest with SMOTE")

```

No handles with labels found to put in legend.






In []:

```
# Random Forest model with downsampling.
rf_d = get_rf_downsampler(x_train, y_train.values.ravel())
predictions = rf_d.predict(x_test)
probabilities = rf_d.predict_proba(x_test)
final_result_upselling["rf_down"] = get_metrics(predictions, y_test.values.ravel(), probabilities)

conf_matrix_upselling["rf_down"] = get_confusion_matrix(y_test, predictions, rf_d.classes_)
```




In []:

```
## ROC Curve for Random Forest model with downsampling.
roc_upselling["rf_down"] = get_roc_curve(rf_d, x_test, y_test, "Random Forest with Downsampling")
```

No handles with labels found to put in legend.

ROC curve for Random Forest with Downsampling.




In []:

```
## Adaboost model
ada, gs_ada = get_adaboost(x_train, y_train.values.ravel())
predictions = ada.predict(x_test)
probabilities = ada.predict_proba(x_test)
final_result_upselling["adaboost"] = get_metrics(predictions, y_test.values.ravel(), probabilities)

conf_matrix_upselling["adaboost"] = get_confusion_matrix(y_test, predictions, ada.classes_)

{'algorithm': 'SAMME', 'base_estimator_ccp_alpha': 0.0, 'base_estimator_class_weight': None, 'base_estimator_criterion': 'gini', 'base_estimator_max_depth': 1, 'base_estimator_max_features': None, 'base_estimator_max_leaf_nodes': None, 'base_estimator_min_impurity_decrease': 0.0, 'base_estimator_min_samples_leaf': 1, 'base_estimator_min_samples_split': 2, 'base_estimator_min_weight_fraction_leaf': 0.0, 'base_estimator_random_state': None, 'base_estimator_splitter': 'best', 'base_estimator': DecisionTreeClassifier(max_depth=1), 'learning_rate': 1, 'n_estimators': 100, 'random_state': None}
```






In []:

```
## ROC Curve for adaboost model.
roc_upselling["adaboost"] = get_roc_curve(ada, x_test, y_test, "Adaboost")
```


No handles with labels found to put in legend.



In []:

```
## Printing the comparison metrics.
print_metrics(final_result_upselling)
```

Model Name	Accuracy	Recall	Precision	F1
Score	AUC score			



In []:


```
plt.rcParams["figure.figsize"] = (5,5)
sns.boxplot(feat_importances)
```

/usr/local/lib/python3.7/dist-packages/seaborn/_decorators.py:43: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

FutureWarning

Out[]:

<matplotlib.axes._subplots.AxesSubplot at 0x7fe7ac8ad2d0>





In []:

```
## Finding important features using Random Forest Feature Selection.
rf_feat_select = RandomForestClassifier(n_estimators=340)
rf_feat_select.fit(x_train, y_train.values.ravel())
rf_imp = rf_feat_select.feature_importances_
rf_importance = pd.Series(rf_imp, df_train.columns)

# Select all features above the 25th percentile
ft_rf = rf_importance[rf_importance >= rf_importance.quantile(0.25)].index.tolist()

plt.rcParams["figure.figsize"] = (20,20)
rf_importance.plot(kind= "barh", color = "teal")
plt.show()
```





In []:

```
## Get the final list of important features based on Information Gain and RF importance
final_features = get_feature_intersection(ft_ig, ft_rf)
```

In []:

```
## Reset the data frames based on the new final features
x_train = pd.DataFrame(x_train, columns=df_train.columns)
x_test = pd.DataFrame(x_test, columns=df_train.columns)

x_train = x_train[final_features].values
x_test = x_test[final_features].values

final_result_appetency = {}
conf_matrix_appetency = {}
roc_appetency = {}
```


In []:

```
## Logistic Regression
lreg, gs_lreg = get_logistic_regression(x_train, y_train.values.ravel())
predictions = lreg.predict(x_test)
probabilities = lreg.predict_proba(x_test)
final_result_appetency["logic_reg"] = get_metrics(predictions, y_test.values.ravel(), probabilities)

conf_matrix_appetency["logic_reg"] = get_confusion_matrix(y_test, predictions, lreg.classes_)

{'C': 0.001, 'class_weight': 'balanced', 'dual': False, 'fit_intercept': True, 'intercept_scaling': 1, 'l1_ratio': None, 'max_iter': 5000, 'multi_class': 'auto', 'n_jobs': -1, 'penalty': 'l2', 'random_state': 42, 'solver': 'lbfgs', 'tol': 0.0001, 'verbose': 0, 'warm_
```


```
start': False}
```




```
In [ ]:
```

```
## ROC Curve  
roc_appetency["logic_reg"] = get_roc_curve(lreg, x_test, y_test, "Logistic Regression")
```

No handles with labels found to put in legend.






In []:

```
## Decision Tree
dst, gs_dst = get_decision_tree(x_train, y_train.values.ravel())
predictions = dst.predict(x_test)
probabilities = dst.predict_proba(x_test)
final_result_appetency["decision_tree"] = get_metrics(predictions, y_test.values.ravel(), probabilities)

conf_matrix_appetency["decision_tree"] = get_confusion_matrix(y_test, predictions, dst.classes_)


{'ccp_alpha': 0.0, 'class_weight': None, 'criterion': 'gini', 'max_depth': 6, 'max_features': None, 'max_leaf_nodes': None, 'min_impurity_decrease': 0.0, 'min_samples_leaf': 1, 'min_samples_split': 3, 'min_weight_fraction_leaf': 0.0, 'random_state': 42, 'splitter': 'best'}
```




In []:

```
## ROC Curve
roc_appetency["decision_tree"] = get_roc_curve(dst, x_test, y_test, "Decision Tree")
```

No handles with labels found to put in legend.





In []:


```
## Random Vanilla Forest
rfv, gs_rfv = get_rf_vanilla(x_train, y_train.values.ravel())
predictions = rfv.predict(x_test)
probabilities = rfv.predict_proba(x_test)
final_result_appetency["vanilla_random_forest"] = get_metrics(predictions, y_test.values
.ravel(), probabilities)

conf_matrix_appetency["vanilla_random_forest"] = get_confusion_matrix(y_test, predictions
, rfv.classes_)
```

```
{'bootstrap': True, 'ccp_alpha': 0.0, 'class_weight': None, 'criterion': 'gini', 'max_dep
th': 10, 'max_features': 'sqrt', 'max_leaf_nodes': None, 'max_samples': None, 'min_impuri
ty_decrease': 0.0, 'min_samples_leaf': 1, 'min_samples_split': 2, 'min_weight_fraction_le
af': 0.0, 'n_estimators': 80, 'n_jobs': -1, 'oob_score': True, 'random_state': 0, 'verbos
e': 0, 'warm_start': False}
```

```
/usr/local/lib/python3.7/dist-packages/sklearn/metrics/_classification.py:1318: Undefined
MetricWarning: Precision is ill-defined and being set to 0.0 due to no predicted samples.
Use `zero_division` parameter to control this behavior.
_warn_prf(average, modifier, msg_start, len(result))
```






In []:

```
## ROC Curve
roc_appetency["vanilla_random_forest"] = get_roc_curve(rfv, x_test, y_test, "Vanilla Random Forest")
```

No handles with labels found to put in legend.




In []:

```
## Random Forest with Balancing
rf_balance = get_rf_balance(x_train, y_train.values.ravel())
predictions = rf_balance.predict(x_test)
probabilities = rf_balance.predict_proba(x_test)
final_result_appetency["rf_balanced"] = get_metrics(predictions, y_test.values.ravel(),
probabilities)

conf_matrix_appetency["rf_balanced"] = get_confusion_matrix(y_test, predictions, rf_balanc
```

```
ce.classes_)
```





```
In [ ]:
```

```
## ROC Curve  
roc_appetency["rf_balanced"] = get_roc_curve(rf_balance, x_test, y_test, "Balanced Random  
Forest")
```

No handles with labels found to put in legend.

ROC curve for Balanced Random Forest.






In []:

```
## Random Forest with SMOTE resampling
rf_smote = get_rf_smote(x_train, y_train.values.ravel())
predictions = rf_smote.predict(x_test)
probabilities = rf_smote.predict_proba(x_test)
final_result_appetency["rf_smote"] = get_metrics(predictions, y_test.values.ravel(), probabilities)


conf_matrix_appetency["rf_smote"] = get_confusion_matrix(y_test, predictions, rf_smote.classes_)
```




In []:

```
## ROC Curve
roc_appetency["rf_smote"] = get_roc_curve(rf_smote, x_test, y_test, "Random Forest with SMOTE")
```

No handles with labels found to put in legend.






In []:

```
## Random Forest with downsampling
rf_d = get_rf_downsampler(x_train, y_train.values.ravel())
predictions = rf_d.predict(x_test)
probabilities = rf_d.predict_proba(x_test)
final_result_appetency["rf_down"] = get_metrics(predictions, y_test.values.ravel(), probabilities)


conf_matrix_appetency["rf_down"] = get_confusion_matrix(y_test, predictions, rf_d.classes_)
```



In []:

```
## ROC Curve
roc_appetency["rf_down"] = get_roc_curve(rf_d, x_test, y_test, "Random Forest with SMOTE")
```

No handles with labels found to put in legend.



In []:


```
## AdaBoost
ada, gs_ada = get_adaboost(x_train, y_train.values.ravel(), 'RECALL')
predictions = ada.predict(x_test)
probabilities = ada.predict_proba(x_test)
final_result_appetency["adaboost"] = get_metrics(predictions, y_test.values.ravel(), probabilities)
```

```
conf_matrix_appetency["adaboost"] = get_confusion_matrix(y_test, predictions, ada.classes_)
```

```
{'algorithm': 'SAMME', 'base_estimator_ccp_alpha': 0.0, 'base_estimator_class_weight': None, 'base_estimator_criterion': 'gini', 'base_estimator_max_depth': 1, 'base_estimator_max_features': None, 'base_estimator_max_leaf_nodes': None, 'base_estimator_min_impurity_decrease': 0.0, 'base_estimator_min_samples_leaf': 1, 'base_estimator_min_samples_split': 2, 'base_estimator_min_weight_fraction_leaf': 0.0, 'base_estimator_random_state': None, 'base_estimator_splitter': 'best', 'base_estimator': DecisionTreeClassifier(max_depth=1), 'learning_rate': 0.01, 'n_estimators': 20, 'random_state': None}
```

/usr/local/lib/python3.7/dist-packages/sklearn/metrics/_classification.py:1318: Undefined MetricWarning: Precision is ill-defined and being set to 0.0 due to no predicted samples. Use `zero_division` parameter to control this behavior.


```
_warn_prf(average, modifier, msg_start, len(result))
```




In []:

```
## ROC Curve
roc_appetency["adaboost"] = get_roc_curve(ada, x_test, y_test, "Adaboost")
```

No handles with labels found to put in legend.





In []:

```
print_metrics(final_result_appetency)
```

Model Name		Accuracy	Recall	Precision	
F1 Score	AUC score				
logic_reg	0.7305	0.7696629213483146	0.049086348978860626	0.0922	
8696530818457	0.8187166965239694				
decision_tree	0.9821	0.0	0.0		
0.0	0.8177511960080444				
vanilla_random_forest	0.9822	0.0	0.0		
0.0	0.810392972437477				
rf_balanced	0.9008	0.4438202247191011	0.08127572016460906	0.1373	
9130434782612	0.8008191882931919				
rf_smote	0.8973	0.42134831460674155	0.07507507507507508	0.1274	
4265080713676	0.7805185103837065				
rf_down	0.7144	0.7528089887640449	0.04548540393754243	0.0857	
8745198463508	0.800674763601088				
adaboost	0.9822	0.0	0.0		
0.0	0.5				

In []:

```
## The permutation importance for rf_down
get_permutation_importance(rf_balance, x_test, y_test, final_features)
```

Var208_kIsH0.000 +/- 0.000

Predicting Churn

In []:

```
## Splitting the training data as there is no corresponding label file
## for orange_small_test.data
## Get the test train split with Upselling data
x_train, x_test, y_train, y_test = get_train_test_split(df_train, y_train_churn)
```


In []:

```
## Compute information gain for features
ig_importances = mutual_info_classif(x_train, y_train.values.ravel())
feat_importances = pd.Series(ig_importances, df_train.columns)
ft_ig = feat_importances[feat_importances >= feat_importances.quantile(0.25)].index.tolist()

plt.rcParams["figure.figsize"] = (20,20)
feat_importances.plot(kind= "barh", color = "teal")

plt.show()
```





In []:


```
plt.rcParams["figure.figsize"] = (5,5)
sns.boxplot(feat_importances)
```

/usr/local/lib/python3.7/dist-packages/seaborn/_decorators.py:43: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

FutureWarning

Out[]:

<matplotlib.axes._subplots.AxesSubplot at 0x7fe7a45d4a90>




In []:

```
## Random Forest Feature Importance
rf_feat_select = RandomForestClassifier(n_estimators=340)
```

```
rf_feat_select.fit(x_train, y_train.values.ravel())
rf_imp = rf_feat_select.feature_importances_
rf_importance = pd.Series(rf_imp, df_train.columns)

# Select all features above the 25th percentile
ft_rf = rf_importance[rf_importance >= rf_importance.quantile(0.25)].index.tolist()

plt.rcParams["figure.figsize"] = (20,20)
rf_importance.plot(kind= "barh", color = "teal")
plt.show()
```



In []:

```
## Get the final list of important features based on Information Gain and RF importance
final_features = get_feature_intersection(ft_ig, ft_rf)
print(final_features)
```

```
['Var6', 'Var7', 'Var13', 'Var21', 'Var22', 'Var24', 'Var25', 'Var28', 'Var35', 'Var38',
'Var57', 'Var65', 'Var73', 'Var74', 'Var76', 'Var78', 'Var81', 'Var83', 'Var85', 'Var112',
, 'Var113', 'Var119', 'Var123', 'Var125', 'Var126', 'Var132', 'Var133', 'Var134', 'Var140',
, 'Var143', 'Var144', 'Var153', 'Var160', 'Var163', 'Var181', 'Var203_9_Y1', 'Var205_09_Q',
'Var205_VpdQ', 'Var205_sJzTlal', 'Var208_kIsH', 'Var210_uKAI', 'Var211_L84s', 'Var218_UYBR',
'Var218_cJvF', 'Var221_oslk', 'Var221_zCkv', 'Var223_LM8l689qOp', 'Var227_02N6s8f',
, 'Var227_RAYp', 'Var227_ZI9m']
```

In []:

```
## Reset the data frames based on the new final features
x_train = pd.DataFrame(x_train, columns=df_train.columns)
```

```
x_test = pd.DataFrame(x_test, columns=df_train.columns)
```

```
x_train = x_train[final_features].values  
x_test = x_test[final_features].values
```


```
final_result_churn = {}  
conf_matrix_churn = {}  
roc_churn = {}
```

In []:

```
## Logistic Regression
```

```
lreg, gs_lreg = get_logistic_regression(x_train, y_train.values.ravel())  
predictions = lreg.predict(x_test)  
probabilities = lreg.predict_proba(x_test)  
final_result_churn["logic_reg"] = get_metrics(predictions, y_test.values.ravel(), probabilities)  
  
conf_matrix_churn["logic_reg"] = get_confusion_matrix(y_test, predictions, lreg.classes_)
```

```
{'C': 0.1, 'class_weight': 'balanced', 'dual': False, 'fit_intercept': True, 'intercept_scaling': 1, 'l1_ratio': None, 'max_iter': 5000, 'multi_class': 'auto', 'n_jobs': -1, 'penalty': 'l2', 'random_state': 42, 'solver': 'lbfgs', 'tol': 0.0001, 'verbose': 0, 'warm_start': False}
```

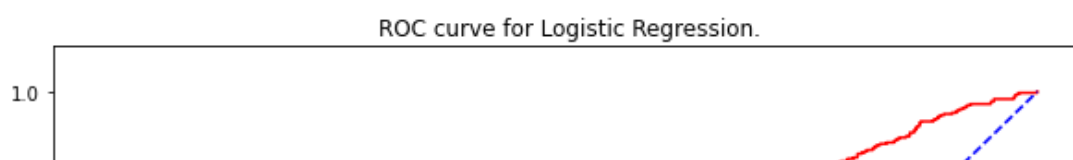



In []:

```
## ROC Curve
```

```
roc_churn["logic_reg"] = get_roc_curve(lreg, x_test, y_test, "Logistic Regression")
```

No handles with labels found to put in legend.






In []:

```
## Decision Tree
dst, gs_dst = get_decision_tree(x_train, y_train.values.ravel())
predictions = dst.predict(x_test)
probabilities = dst.predict_proba(x_test)
final_result_churn["decision_tree"] = get_metrics(predictions, y_test.values.ravel(), probabilities)

conf_matrix_churn["decision_tree"] = get_confusion_matrix(y_test, predictions, dst.classes_)

{'ccp_alpha': 0.0, 'class_weight': None, 'criterion': 'gini', 'max_depth': 6, 'max_features': None, 'max_leaf_nodes': None, 'min_impurity_decrease': 0.0, 'min_samples_leaf': 1, 'min_samples_split': 3, 'min_weight_fraction_leaf': 0.0, 'random_state': 42, 'splitter': 'best'}
```






In []:

```
## ROC Curve
roc_churn["decision_tree"] = get_roc_curve(dst, x_test, y_test, "Decision Tree")
```

No handles with labels found to put in legend.




In []:

```
## Random Vanilla Forest
rfv, gs_rfv = get_rf_vanilla(x_train, y_train.values.ravel())
predictions = rfv.predict(x_test)
probabilities = rfv.predict_proba(x_test)
final_result_churn["vanilla_random_forest"] = get_metrics(predictions, y_test.values.ravel(), probabilities)

conf_matrix_churn["vanilla_random_forest"] = get_confusion_matrix(y_test, predictions, rfv.classes_)
```

```
{'bootstrap': True, 'ccp_alpha': 0.0, 'class_weight': None, 'criterion': 'gini', 'max_depth': 10, 'max_features': 'sqrt', 'max_leaf_nodes': None, 'max_samples': None, 'min_impurity_decrease': 0.0, 'min_samples_leaf': 1, 'min_samples_split': 2, 'min_weight_fraction_leaf': 0.0, 'n_estimators': 100, 'n_jobs': -1, 'oob_score': True, 'random_state': 0, 'verbose': 0, 'warm_start': False}
```


```
/usr/local/lib/python3.7/dist-packages/sklearn/metrics/_classification.py:1318: Undefined
MetricWarning: Precision is ill-defined and being set to 0.0 due to no predicted samples.
Use `zero_division` parameter to control this behavior.
_warn_prf(average, modifier, msg_start, len(result))
```




In []:

```
## ROC Curve
roc_churn["vanilla_random_forest"] = get_roc_curve(rfv, x_test, y_test, "Vanilla Random F
orest")
```

No handles with labels found to put in legend.






In []:

```
## Random Forest with Balancing
rf_balance = get_rf_balance(x_train, y_train.values.ravel())
predictions = rf_balance.predict(x_test)
probabilities = rf_balance.predict_proba(x_test)
final_result_churn["rf_balanced"] = get_metrics(predictions, y_test.values.ravel(), probabilities)


conf_matrix_churn["rf_balanced"] = get_confusion_matrix(y_test, predictions, rf_balance.classes_)
```




In []:

```
## ROC Curve
roc_churn["rf_balanced"] = get_roc_curve(rf_balance, x_test, y_test, "Balanced Random Forest")
```

No handles with labels found to put in legend.






In []:

```
## Random Forest with SMOTE resampling
rf_smote = get_rf_smote(x_train, y_train.values.ravel())
predictions = rf_smote.predict(x_test)
probabilities = rf_smote.predict_proba(x_test)
final_result_churn["rf_smote"] = get_metrics(predictions, y_test.values.ravel(), probabilities)

conf_matrix_churn["rf_smote"] = get_confusion_matrix(y_test, predictions, rf_smote.classes_)
```






In []:

```
## ROC Curve
roc_churn["rf_smote"] = get_roc_curve(rf_smote, x_test, y_test, "Random Forest with SMOTE")
```

No handles with labels found to put in legend.




In []:

```
## Random Forest with downsampling
rf_d = get_rf_downsampler(x_train, y_train.values.ravel())
predictions = rf_d.predict(x_test)
probabilities = rf_d.predict_proba(x_test)
final_result_churn["rf_down"] = get_metrics(predictions, y_test.values.ravel(), probabilities)

conf_matrix_churn["rf_down"] = get_confusion_matrix(y_test, predictions, rf_d.classes_)
```






In []:

```
## ROC Curve
roc_churn["rf_down"] = get_roc_curve(rf_d, x_test, y_test, "Random Forest with Down Sampling")
```

No handles with labels found to put in legend.



In []:

```
## AdaBoost
```

```

## AdaBoost
ada, gs_ada = get_adaboost(x_train, y_train.values.ravel())
predictions = ada.predict(x_test)
probabilities = ada.predict_proba(x_test)
final_result_churn["adaboost"] = get_metrics(predictions, y_test.values.ravel(), probabilities)

conf_matrix_churn["adaboost"] = get_confusion_matrix(y_test, predictions, ada.classes_)

```

```


{'algorithm': 'SAMME', 'base_estimator_ccp_alpha': 0.0, 'base_estimator_class_weight':
None, 'base_estimator_criterion': 'gini', 'base_estimator_max_depth': 1, 'base_estimator_max_features': None, 'base_estimator_max_leaf_nodes': None, 'base_estimator_min_impurity_decrease': 0.0, 'base_estimator_min_samples_leaf': 1, 'base_estimator_min_samples_split': 2, 'base_estimator_min_weight_fraction_leaf': 0.0, 'base_estimator_random_state': None, 'base_estimator_splitter': 'best', 'base_estimator': DecisionTreeClassifier(max_depth=1), 'learning_rate': 1, 'n_estimators': 40, 'random_state': None}

```

```

/usr/local/lib/python3.7/dist-packages/sklearn/metrics/_classification.py:1318: Undefined
MetricWarning: Precision is ill-defined and being set to 0.0 due to no predicted samples.
Use `zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, msg_start, len(result))

```




In []:


```

## ROC Curve
roc_churn["adaboost"] = get_roc_curve(ada, x_test, y_test, "Adaboost")

```

No handles with labels found to put in legend.





```
In [ ]:
```

```
print_metrics(final_result_churn)
```

Model Name	Accuracy	Recall	Precision	F1 Score	AUC score
logic_reg	0.5937	0.6267029972752044	0.10826076723935044	0.18462773429660848	0.6544302777550695
decision_tree	0.9266	0.01634877384196185	0.5	0.031662269129287594	0.6783559742894095
vanilla_random_forest	0.9266	0.0	0.0	0.0	0.6905298501274179
rf_balanced	0.7567	0.4550408719346049	0.14110688635403465	0.21541438245727187	0.6876650800941709
rf_smote	0.8306	0.28746594005449594	0.15267727930535455	0.19943289224952743	0.6637738625463224
rf_down	0.6199	0.6866485013623979	0.12368098159509203	0.2096069868995633	0.69675532887807
adaboost	0.9266	0.0	0.0	0.0	0.6918601949878581

```
In [ ]:
```

```
## The permutation importance for rf_balance
get_permutation_importance(rf_balance, x_test, y_test, final_features)
```

```
Var126 0.022 +/- 0.002
Var218_UYBR0.002 +/- 0.001
Var143 0.000 +/- 0.000
```